

Yocto Project and OpenEmbedded Training

STM32MP1 variant

Practical Labs


<https://bootlin.com>

July 22, 2020

About this document

Updates to this document can be found on <https://bootlin.com/doc/training/yocto-stm32>.

This document was generated from LaTeX sources found on <https://github.com/bootlin/training-materials>.

More details about our training sessions can be found on <https://bootlin.com/training>.

Copying this document

© 2004-2020, Bootlin, <https://bootlin.com>.



This document is released under the terms of the [Creative Commons CC BY-SA 3.0 license](#). This means that you are free to download, distribute and even modify it, under certain conditions.

Corrections, suggestions, contributions and translations are welcome!

Training setup

Download files and directories used in practical labs

Install lab data

For the different labs in this course, your instructor has prepared a set of data (kernel images, kernel configurations, root filesystems and more). Download and extract its tarball from a terminal:

```
cd
wget https://bootlin.com/doc/training/sessions/online.yocto.jul2020/yocto-stm32-labs.tar.xz
tar xvf yocto-stm32-labs.tar.xz
```

Lab data are now available in an `yocto-stm32-labs` directory in your home directory. This directory contains directories and files used in the various practical labs. It will also be used as working space, in particular to keep generated files separate when needed.

You are now ready to start the real practical labs!

Install extra packages

Feel free to install other packages you may need for your development environment. In particular, we recommend to install your favorite text editor and configure it to your taste. The favorite text editors of embedded Linux developers are of course *Vim* and *Emacs*, but there are also plenty of other possibilities, such as Visual Studio Code¹, *GEdit*, *Qt Creator*, *CodeBlocks*, *Geany*, etc.

It is worth mentioning that by default, Ubuntu comes with a very limited version of the `vi` editor. So if you would like to use `vi`, we recommend to use the more featureful version by installing the `vim` package.

More guidelines

Can be useful throughout any of the labs

- Read instructions and tips carefully. Lots of people make mistakes or waste time because they missed an explanation or a guideline.
- Always read error messages carefully, in particular the first one which is issued. Some people stumble on very simple errors just because they specified a wrong file path and didn't pay enough attention to the corresponding error message.
- Never stay stuck with a strange problem more than 5 minutes. Show your problem to your colleagues or to the instructor.
- You should only use the `root` user for operations that require super-user privileges, such as: mounting a file system, loading a kernel module, changing file ownership, configuring

¹This tool from Microsoft is Open Source! To try it on Ubuntu: `sudo snap install code`

the network. Most regular tasks (such as downloading, extracting sources, compiling...) can be done as a regular user.

- If you ran commands from a root shell by mistake, your regular user may no longer be able to handle the corresponding generated files. In this case, use the `chown -R` command to give the new files back to your regular user.

Example: `chown -R myuser.myuser linux/`

Lab1: First Yocto Project build

Your first dive into Yocto Project and its build mechanism

During this lab, you will:

- Set up an OpenEmbedded environment
- Configure the project and choose a target
- Build your first Poky image

Setup

Before starting this lab, make sure your home directory is not encrypted. OpenEmbedded cannot be used on top of an eCryptFS file system due to limitations in file name lengths.

Go to the `$HOME/yocto-stm32-labs/` directory.

Install the required packages:

```
sudo apt install bc build-essential chrpath cpio diffstat gawk git python texinfo wget gdisk
```

Download Yocto

Download the thud version of Poky:

```
git clone git://git.yoctoproject.org/poky.git
cd $HOME/yocto-stm32-labs/poky
git checkout -b zeus-22.0.3 zeus-22.0.3
```

Return to your project root directory (`cd $HOME/yocto-stm32-labs/`) and download the OpenEmbedded and STM32MP layers:

```
git clone -b zeus git://git.openembedded.org/meta-openembedded
git clone -b zeus https://github.com/STMicroelectronics/meta-st-stm32mp.git
```

Set up the build environment

Check you're using Bash. This is the default shell when using Ubuntu.

Export all needed variables and set up the build directory:

```
cd $HOME/yocto-stm32-labs
source poky/oe-init-build-env
```

You must specify which machine is your target. By default it is `qemu`. We need to build an image for an `stm32mp1`. Update the `MACHINE` configuration variable accordingly.

Also, if you need to save disk space on your computer you can add `INHERIT += "rm_work"` in the previous configuration file. This will remove the package work directory once a package is built.

Don't forget to make the configuration aware of the OpenEmbedded and STM32MP layer. Edit the layer configuration file (`$BUILDDIR/conf/bblayers.conf`) and append the full path to the `meta-openembedded/meta-oe`, `meta-openembedded/meta-python` and `meta-st-stm32mp` directory to the `BBLAYERS` variable.

Build your first image

Now that you're ready to start the compilation, simply run:

```
bitbake core-image-minimal
```

Once the build finished, you will find the output images under `$BUILDDIR/tmp/deploy/images/stm32mp1`.

Set up the SD card

In this first lab we will use an SD card to store the bootloader, kernel and root filesystem files. To generate the final image, You will find scripts in `$BUILDDIR/tmp/deploy/images/stm32mp1/scripts`.

Execute it:

```
sudo ./create_sdcard_from_flashlayout.sh \  
../flashlayout_core-image-minimal/FlashLayout_sdcard_stm32mp157a-dk1-basic.tsv
```

Flash the SD card with that image:

```
umount /dev/mmcblk0p*  
sudo dd if=../flashlayout_core-image-minimal/FlashLayout_sdcard_stm32mp157a-dk1-basic.raw \  
of=/dev/mmcblk0 bs=8M conv=fdatasync status=progress
```

Setting up serial communication with the board

Plug the USB-A to micro USB-B cable on the Discovery board. There is only one micro USB port on the board, it is CN11, also named ST-LINK. This is a debug interface and exposes multiple debugging interfaces, including a serial interface. When plugged in your computer, a serial port should appear, `/dev/ttyACM0`.

You can also see this device appear by looking at the output of `dmesg`.

To communicate with the board through the serial port, install a serial communication program, such as `picocom`:

```
sudo apt install picocom
```

You also need to make your user belong to the `dialout` group to be allowed to write to the serial console:

```
sudo adduser $USER dialout
```

Important: for the group change to be effective, in Ubuntu 18.04, you have to *completely reboot* the system ². A workaround is to run `newgrp dialout`, but it is not global. You have to run it in each terminal.

²As explained on <https://askubuntu.com/questions/1045993/after-adding-a-group-logoutlogin-is-not-enough-in-18-04/>.

Run `picocom -b 115200 /dev/ttyACM0`, to start serial communication on `/dev/ttyACM0`, with a baudrate of 115200.

If you wish to exit `picocom`, press `[Ctrl][a]` followed by `[Ctrl][x]`.

There should be nothing on the serial line so far, as the board is not powered up yet.

Boot

Insert the SD card in the dedicated slot on the Discovery. You can now power-up the board by connecting the USB-C cable to the board, in CN6, PWR_IN and to your PC at the other end. You should see boot messages on the console. Wait until the login prompt, then enter `root` as user. Congratulations! The board has booted and you now have a shell.

Lab2: Advanced Yocto configuration

Configure the build, customize the output images and use NFS

During this lab, you will:

- Customize the package selection
- Configure the build system
- Use the rootfs over NFS

Set up the Ethernet communication and NFS on the board

It isn't practical at all to reflash the root filesystem on the target every time a change is made. Fortunately, it is possible to set up networking between the development workstation and the target. Then, workstation files can be accessed by the target through the network, using NFS.

First we need to set the kernel boot arguments U-Boot will pass to the Linux kernel at boot time. For that, edit the `mmc0_stm32mp157a-dk1_extlinux/extlinux.conf` file in the bootfs partition of the SD card and change the APPEND line for the `stm32mp157a-dk1-sdcard` configuration:

```
APPEND root=/dev/nfs rw console=ttySTM0,115200 nfsroot=192.168.0.1:/nfs,vers=3,tcp ip=192.168.0.100:::et
```

Set up the Ethernet communication on the workstation

With a network cable, connect the Ethernet port of your board to the one of your computer. If your computer already has a wired connection to the network, your instructor will provide you with a USB Ethernet adapter. A new network interface should appear on your Linux system.

Install the `net-tools` package and find the name of this interface by typing:

```
ifconfig -a
```

The network interface name is likely to be `enxxx`³. If you have a pluggable Ethernet device, it's easy to identify as it's the one that shows up after plugging in the device.

Then, instead of configuring the host IP address from NetWork Manager's graphical interface, let's do it through its command line interface, which is so much easier to use:

```
nmcli con add type ethernet ifname en... ip4 192.168.0.1/24
```

Set up the NFS server on the workstation

First install the NFS server on the training computer and create the root NFS directory:

³Following the *Predictable Network Interface Names* convention: <https://www.freedesktop.org/wiki/Software/systemd/PredictableNetworkInterfaceNames/>


```
sudo apt install nfs-kernel-server
sudo mkdir -m 777 /nfs
```

Then make sure this directory is used and exported by the NFS server by adding `/nfs *(rw, sync, no_root_squash, subtree_check)` to the `/etc/exports` file.

Finally restart the service:

```
sudo service nfs-kernel-server restart
```

Add a package to the rootfs image

You can add packages to be built by editing the local configuration file `$BUILDDIR/conf/local.conf`. The `IMAGE_INSTALL` variable controls the packages included into the output image.

To illustrate this, add the Dropbear SSH server to the list of enabled packages.

Tip: do not override the default enabled package list, but append the Dropbear package instead.

Boot with the updated rootfs

First we need to put the rootfs under the NFS root directory so that it is accessible by NFS clients. Simply uncompress the archived output image in the previously created `/nfs` directory:

```
sudo tar xpf $BUILDDIR/tmp/deploy/images/stm32mp1/\
  core-image-minimal-stm32mp1.tar.xz -C /nfs
```

Then boot the board.

The Dropbear SSH server was enabled a few steps before, and should now be running as a service on the Discovery. You can test it by accessing the board through SSH:

```
ssh root@192.168.0.100
```

You should see the Discovery command line!

Choose a package variant

Dependencies of a given package are explicitly defined in its recipe. Some packages may need a specific library or piece of software but others only depend on a functionality. As an example, the kernel dependency is described by `virtual/kernel`.

To see which kernel is used, dry-run BitBake:

```
bitbake -vn virtual/kernel
```

In our case, we can see the `linux-stm32mp` provides the `virtual/kernel` functionality:

```
NOTE: selecting linux-stm32mp to satisfy virtual/kernel due to PREFERRED_PROVIDERS
```

We can force bitbake to select another kernel by explicitly defining which one to use in our local configuration. Try switching from `linux-stm32mp` to `linux-dummy` only using the local configuration.

Then check the previous step worked by dry-running again BitBake.

```
bitbake -vn virtual/kernel
```

You can now rebuild the whole Yocto project, with `bitbake core-image-minimal`

Tip: you need to define the more specific information here to be sure it is the one used. The `MACHINE` variable can help here.

As this was only to show how to select a preferred provider for a given package, you can now use `linux-stm32mp` again.

BitBake tips

BitBake is a powerful tool which can be used to execute specific commands. Here is a list of some useful ones, used with the `virtual/kernel` package.

- The Yocto recipes are divided into numerous tasks, you can print them by using: `bitbake -c listtasks virtual/kernel`.
- BitBake allows to call a specific task only (and its dependencies) with: `bitbake -c <task> virtual/kernel`. (<task> can be `menuconfig` here).
- You can force to rebuild a package by calling: `bitbake -f virtual/kernel`
- `world` is a special keyword for all packages. `bitbake --runall=fetch world` will download all packages sources (and their dependencies).
- You can get a list of locally available packages and their current version with: `bitbake -s`
- You can also find detailed information on available packages, their current version, dependencies or the contact information of the maintainer by visiting:
<http://recipes.yoctoproject.org/>

For detailed information, please run `bitbake -h`

Going further

If you have some time left, let's improve our setup to use TFTP, in order to avoid having to reflash the SD card for every test. What you need to do is:

1. Install a TFTP server (package `tftpd-hpa`) on your system.
2. Copy the Linux kernel image and Device Tree to `/var/lib/tftpboot` so that they are made available by the TFTP server.
3. Change the U-Boot `bootcmd` to load the kernel image and the Device Tree over TFTP.

See the training materials of our *Embedded Linux system development* course for details!

Lab3: Add a custom application

Add a new recipe to support a required custom application

During this lab, you will:

- Write a recipe for a custom application
- Integrate this application into the build

This is the first step of adding an application to Yocto. The remaining part is covered in the next lab, "Create a Yocto layer".

Setup and organization

In this lab we will add a recipe handling the `nInvaders` application. Before starting the recipe itself, find the `recipes-extended` directory in the OpenEmbedded-Core layer and add a subdirectory for your application.

A recipe for an application is usually divided into a version specific `bb` file and a common one. Try to follow this logic and separate the configuration variables accordingly.

Tip: it is possible to include a file into a recipe with the keyword `require`.

First hands on nInvaders

The `nInvaders` application is a terminal based game following the space invaders family. In order to deal with the text based user interface, `nInvaders` uses the `ncurses` library.

First try to find the project homepage, download the sources and have a first look: license, Makefile, requirements...

Write the common recipe

Create an appropriate common file, ending in `.inc`

In this file add the common configuration variables: source URI, description...

Write the version specific recipe

Create a file that respects the Yocto nomenclature: `${PN}_${PV}.bb`

Add the required common configuration variables: archive checksum, license file checksum, package revision...

Testing and troubleshooting

You can check the whole packaging process is working fine by explicitly running the build task on the `nInvaders` recipe:

bitbake ninvaders

Try to make the recipe on your own. Also eliminate the warnings related to your recipe: some configuration variables are not mandatory but it is a very good practice to define them all.

If you hang on a problem, check the following points:

- The common recipe is included in the version specific one
- The checksum and the URI are valid
- The dependencies are explicitly defined
- The internal state has changed, clean the working directory:

```
bitbake -c cleanall ninvaders
```

Tip: BitBake has command line flags to increase its verbosity and activate debug outputs. Also, remember that you need to cross-compile nInvaders for ARM ! Maybe, you will have to configure your recipe to resolve some mistakes done in the application's Makefile (which is often the case). A bitbake variable permits to add some Makefile's options, you should look for it.

Update the rootfs and test

Now that you've compiled the nInvaders application, generate a new rootfs image with `bitbake core-image-minimal`. Then update the NFS root directory. You can confirm the nInvaders program is present by running:

```
find /nfs -iname ninvaders
```

Access the board command line through SSH. You should be able to launch the nInvaders program. Now, it's time to play!

Inspect the build

The nInvaders application was unpacked and compiled in the recipe's work directory. Can you spot nInvaders' directory in the build work directory?

Once you found it, look around. You should at least spot some directories:

- The sources. Remember the `${S}` variable?
- `temp`. There are two kinds of files in there. Can you tell what are their purposes?
- Try to see if the licences of nInvaders were extracted.

Lab4: Create a Yocto layer

Add a custom layer to the Yocto project for your project needs

During this lab, you will:

- Create a new Yocto layer
- Interface this custom layer to the existing Yocto project
- Use applications from custom layers

This lab extends the previous one, in order to fully understand how to interface a custom project to the basic Yocto project.

Tools

You can access the configuration and state of layers with the `bitbake-layers` command. This command can also be used to retrieve useful information about available recipes. Try the following commands:

```
bitbake-layers show-layers
bitbake-layers show-recipes linux-yocto
bitbake-layers show-overlayed
bitbake-layers create-layer
```

Create a new layer

With the above commands, create a new Yocto layer named `meta-bootlinlabs` with a priority of 7.

Before using the new layer, we need to configure its generated configuration files. You can start with the `README` file which is not used in the build process but contains information related to layer maintenance. You can then check, and adapt if needed, the global layer configuration file located in the `conf` directory of your custom layer.

Integrate a layer to the build

To be fair, we already used and integrated a layer in our build configuration during the first lab, with `meta-st-stm32`. This layer was responsible for `stm32mp1 DK1` support in Yocto. We have to do the same for our `meta-bootlinlabs` now.

There is a file which contains all the paths of the layers we use. Try to find it without looking back to the first lab. Then add the full path to our newly created layer to the list of layers.

Validate the integration of the `meta-bootlinlabs` layer with:

```
bitbake-layers show-layers
```

and make sure you don't have any warning from bitbake.

Add a recipe to the layer

In the previous lab we introduced a recipe for the `nInvaders` game. We included it to the existing `meta` layer. While this approach give a working result, the Yocto logic is not respected. You should instead always **use a custom layer** to add recipes or to customize the existing ones. To illustrate this we will move our previously created `nInvaders` recipe into the `meta-bootlinlabs` layer.

You can check the `nInvaders` recipe is part of the `meta` layer first:

```
bitbake-layers show-recipes ninvaders
```

Then move the `nInvaders` recipe to the `meta-bootlinlabs` layer. You can check that the `nInvaders` recipe is now part of the layer with the `bitbake-layers` command.

Lab5: Extend a recipe

Add your features to an existing recipe

During this lab, you will:

- Apply patches to an existing recipe
- Use a custom configuration file for an existing recipe
- Extend a recipe to fit your needs

Create a basic appended recipe

To avoid rewriting recipes when a modification is needed on an already existing one, BitBake allows to extend recipes and to overwrite, append or prepend configuration variables values through the so-called BitBake append files.

We will first create a basic BitBake append file, without any change made to the original recipe, to see how it is integrated into the build. We will then extend some configuration variables of the original recipe.

Try to create an appended recipe with the help of the online Yocto Project development documentation. You can find it at <https://www.yoctoproject.org/docs/2.5/dev-manual/dev-manual.html>. We here aim to extend the `linux-stm32mp` kernel recipe.

You can see available `bbappend` files and the recipe they apply to by using the `bitbake-layers` tool (again!):

```
bitbake-layers show-appends
```

If the BitBake append file you just created is recognized by your Yocto environment, you should see:

```
linux-stm32mp_4.19.bb:  
$HOME/yocto-stm32-labs/meta-bootlinlabs/recipes-kernel/linux/linux-stm32mp_4.19.bbappend
```

Add patches to apply in the recipe

We want our extended `linux-stm32mp` kernel to support the Nunchuk as a joystick input. We can add this by applying patches during the `do_patch` task. The needed patches are provided with this lab. You can find them under `~/yocto-stm32-labs/bootlin-lab-data/nunchuk/linux`. For more details about how to write the driver handling the Nunchuk, have a look at our embedded Linux kernel and driver development training course at <https://bootlin.com/training/kernel/>.

Applying a patch is a common task in the daily Yocto process. Many recipes, appended or not, apply a specific patch on top of a mainline project. It's why patches do not have to be explicitly applied, if the recipe inherits from the patch class (directly or not), but only have to be present in the source files list.

Try adding the patches included in this lab to your BitBake append file. Do not forget to also add the `defconfig` file provided alongside the patches. This file contains the kernel configuration. Add `KERNEL_DEFCONFIG = ""` and `KERNEL_EXTERNAL_DEFCONFIG = "defconfig"` so it is used by the `linux-stm32mp` recipe.

You can now rebuild the kernel to take the new patches into account:


```
bitbake virtual/kernel
```

Connect the Nunchuk

Take the Nunchuk device provided by your instructor.

We will connect it to the second I2C port of the CPU (`i2c1`), with pins available on the P9 connector.


Identify the 4 pins of the Nunchuk connector:



Nunchuk i2c pinout
(UEXT connector from Olimex, front view)

Connect the Nunchuk pins:

- The GND pin to CN16 pin 6 (GND)
- The PWR pin to CN16 pin 4 (3V3)
- The DATA pin to CN13 pin 9 (D14)
- The CLK pin to CN13 pin 10 (D15)



Test the Nunchuk

Copy the newly generated kernel and device tree images into the first SD card partition. Then boot the board and wait until you have access to the `busybox` command line.

You can then make sure that the Nunchuk is recognized and is working by checking the presence of the `js0` device file:

```
ls /dev/input/js0
```

Now display the raw events generated by the Nunchuk:

```
cat /dev/input/js0
```

You should see random characters appearing while playing with the Nunchuk. Be aware that the driver we integrated also handles accelerometer events. Therefore, moving the device will produce many events!

Patch nInvaders

The `nInvaders` game uses keyboard events for its controls. We first need to apply a patch introducing joystick support. The patch is located at `~/yocto-stm32-labs/bootlin-lab-data/nunchuk/ninvaders/`.

Add the patch to the `nInvaders SRC_URI`.

Then build a full `core-image-minimal` and update the NFS root directory.

Play nInvaders!

After booting the board you should be able to play `nInvaders` with the keyboard...and the Nunchuk! The `C` button is used to confirm and to fire, and `Z` to pause the game.

Access the board command line through SSH, and launch the game:

```
$ ninvaders
```

Lab6: Create a custom machine configuration

Let Poky know about your hardware!

During this lab, you will:

- Create a custom machine configuration
- Understand how the target architecture is dynamically chosen

Create a custom machine

The machine file configures various hardware related settings. That's what we did in lab1, when we chose the `stm32mp1` one. While it is not necessary to make our custom machine image here, we'll create a new one to demonstrate the process.

Add a new `bootlinlabs` machine to the previously created layer, which will make the Discovery properly boot.

This machine describes a board using the `cortexa7thf-neon-vfpv4` tune and is a part of the `stm32mp` SoC family. Add the following lines to your machine configuration file:

```
require conf/machine/include/st-machine-common-stm32mp.inc
```

```
DEFAULTTUNE = "cortexa7thf-neon-vfpv4"  
require conf/machine/include/tune-cortexa7.inc
```

Populate the machine configuration

This `bootlinlabs` machine needs:

- To define a few variables to set to get the tooling from ST Micro to work properly:

```
BOOTSCHHEME_LABELS += "trusted"  
STM32MP_DT_FILES_DK += "stm32mp157a-dk1"  
ST_KERNEL_LOADADDR = "0xC2000040"  
FLASHLAYOUT_CONFIG_LABELS += "sdcard"
```

- And to support some features:

- `kernel26`
- `apm`
- `usb gadget`
- `usb host`
- `vfat`
- `ext2`

– alsa

Build an image with the new machine

You can now update the MACHINE variable value in the local configuration and start a fresh build.

Check generated files are here and correct

Once the generated images supporting the new `bootlinlabs` machine are generated, you can check all the needed images were generated correctly.

Have a look in the output directory, in `$BUILDDIR/tmp/deploy/images/bootlinlabs/`. Is there something missing?

Update the rootfs

You can now update your root filesystem, to use the newly generated image supporting our `bootlinlabs` machine!

Lab7: Create a custom image

The highest level of customization in Poky

During this lab, you will:

- Write a full customized image recipe
- Choose the exact packages you want on your board

Add a basic image recipe

A build is configured by two top level recipes: the machine recipe and the image one. The image recipe is the top configuration file for the generated rootfs and the packages it includes. Our aim in this lab is to define a custom image from scratch to allow a precise selection of packages on the BeagleBone Black. To show how to deal with real world configuration and how the Yocto Project can be used in the industry we will, in addition to the production image recipe you will use in the final product, create a development one including debug tools and show how to link the two of them to avoid configuration duplication.

First add a custom image recipe in the `meta-bootlinlabs` layer. We will name it `bootlinlabs-image-minimal`. You can find information on how to create a custom image on the dedicated Yocto Project development manual at <https://www.yoctoproject.org/docs/2.5/dev-manual/dev-manual.html>. There are different ways to customize an image, we here want to create a full recipe, using a custom `.bb` file.

Do not forget to inherit from the `core-image` class.

Select the images capabilities and packages

You can control the packages built and included into the final image with the `IMAGE_INSTALL` configuration variable. It is a list of packages to be built. You can also use package groups to include a bunch of programs, generally enabling a functionality, such as `packagegroup-core-boot` which adds the minimal set of packages required to boot an image (i.e. a shell or a kernel).

You can find the package groups under the `packagegroups` directories. To have a list of the available one:

```
find -name packagegroups
```

Open some of them to read their description and have an idea about the capabilities they provide. Then update the installed packages of the image recipe and don't forget to add the `nInvaders` one!

Add a custom package group

We just saw it is possible to use package groups to organize and select the packages instead of having a big blob of configuration in the image recipe itself. We will here create a custom package for game related recipes.

With the above documentation, create a `packagegroup-bootlinlabs-games` group which inherits from the `packagegroup` class. Add the `nInvaders` program into its runtime dependencies.

Now update the image recipe to include the package group instead of the `nInvaders` program directly.

Differentiate the production recipe from the debug one

You can enable the debugging capabilities of your image just by changing the BitBake target when building the whole system. We want here to have a common base for both the production and the debug images, but also take into account the possible differences. In our example only the built package list will change.

Create a debug version of the previous image recipe, and name it `bootlinlabs-image-minimal-dbg`. Try to avoid duplicating code! Then add the `dbg-pkgs` to the image features list. It is also recommended to update the recipe's description, and to add extra debugging tools.

Build the new debug image with BitBake and check the previously included packages are present in the newly generated rootfs.

Lab8: Develop your application in the Poky SDK

Generate and use the Poky SDK

During this lab, you will:

- Build the Poky SDK
- Install the SDK
- Compile an application for the BeagleBone in the SDK environment

Build the SDK

Two SDKs are available, one only embedding a toolchain and the other one allowing for application development. We will use the latter one here.

First, build an SDK for the `bootlinlabs-image-minimal` image, with the `populate_sdk` task.

Once the SDK is generated, a script will be available at `tmp/deploy/sdk`.

Install the SDK

Open a new console to be sure that no extra environment variable is set. We mean to show you how the SDK sets up a fully working environment.

Install the SDK in `$HOME/yocto-stm32-labs/sdk` by executing the script generated at the previous step.

```
$BUILDDIR/tmp/deploy/sdk/poky-glibc-x86_64-bootlinlabs-image-minimal-cortexa8hf-neon-toolchain-2.5.sh
```

Set up the environment

Go into the directory where you installed the SDK (`$HOME/yocto-stm32-labs/sdk`). Source the environment script:

```
source environment-setup-cortexa8hf-vfp-neon-poky-linux-gnueabi
```

Have a look at the exported environment variables:

```
env
```

Compile an application in the SDK

Download the essential Ctrix sources at <https://download.mobatek.net/sources/ctris-0.42-1-src.tar.bz2>

Extract the source in the SDK:

```
tar xf ctris-0.42-1-src.tar.bz2
tar xf ctris-0.42.tar.bz2
cd ctris-0.42
```

Then modify the Makefile, to make sure that the environment variables exported by the SDK script are not overridden.

Compile the application. You can check the application was successfully compiled for the right target by using the `file` command. The `ctris` binary should be an ELF 32-bit LSB executable compiled for ARM.

Finally, you can copy the binary to the board, by using the `scp` command. Then run it and play a bit to ensure it is working fine!